

MLMI17 Advanced Computer Vision Coursework 2

Instructor: Elliott Wu, sw2181@cam.ac.uk

Due: Tuesday, 24 March 2026, noon 12:00PM

Submission Guidelines

Your answers should contain an explanation of what you do. You must also give an interpretation of what the results and visualizations you provide mean – why are the results the way they are? Each question should be labelled and answered separately and distinctly. Total combined length of answers must not exceed 5 sides of A4 (excluding the cover page), minimum 11pt font, 1 inch margins. Please submit the PDF and all your code in a ZIP file. Please use the code base provided in this GitHub repository to run the experiments: <https://github.com/CambridgeCVCourses/CW2>.

1 Task 1 – Camera Pose Estimation and Sparse Reconstruction

1.1 Using COLMAP [3, 4] [20%]

Capture a video of a scene of your choice and run COLMAP using `ns-process-data` to estimate camera poses and reconstruct a sparse point cloud.

- You must identify one successful capture and one failure capture. A failure case is defined as either unsuccessful registration or fewer than 5% of frames being registered.
- For the successful case, visualise the estimated camera poses and sparse point cloud using `Viser`, and include at least one visualisation in your report.
- For the failure case, provide a explanation of why you think the reconstruction failed.

Detailed instructions for running COLMAP and visualising results are available in the GitHub repository. Two example videos are provided for reference, and COLMAP should succeed on both.

1.2 Using π^3 [5] [20%]

Next, use π^3 to estimate camera poses and sparse point clouds for the same scenes you captured. An incomplete script is provided that runs π^3 inference and exports the results in the same format as the outputs of `ns-process-data`.

- You are required to implement the transformation that converts π^3 predictions into the coordinate system convention used by `nerfstudio`.
- Evaluate whether π^3 works on your videos. Compare the results produced by π^3 with those from COLMAP, and summarise the advantages and limitations of the two methods.

2 Task 2 – Novel View Synthesis

2.1 Comparing NeRF [2] and 3DGS [1] [20%]

- Train a NeRF model on the provided `static_scene` using the default `nerfacto` model in `nerfstudio` and the camera poses computed by π^3 .
- Train a 3D Gaussian Splatting (3DGS) model on the same `static_scene` using the default `splatfacto` model in `nerfstudio` and the camera poses computed by π^3 .

- Compare the rendered 3D scenes and explain the differences you observe, such as visual quality, rendering characteristics, and training behaviour. Include visualisations and interpretation of your findings in the report.

2.2 Comparing Static and Dynamic Scenes [10%]

- Train a 3DGS model on the provided `dynamic_scene` using the same default `splatfacto` model in `nerfstudio` and the camera poses computed by π^3 .
- Compare the results obtained for the static and dynamic scenes. Describe and interpret the differences you observe. Include visualisations and interpretation of your findings in the report.

2.3 Improving the Dynamic Scene [30%]

Propose methods to improve the reconstruction and rendering quality of the dynamic scene. You may describe potential solutions even if you are unable to validate them through experiments. Clearly justify why your proposed approaches could address the observed limitations.

References

- [1] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkühler, and George Drettakis. 3d gaussian splatting for real-time radiance field rendering. *TOG*, 42(4), 2023. 1
- [2] Ben Mildenhall, Pratul P. Srinivasan, Matthew Tancik, Jonathan T. Barron, Ravi Ramamoorthi, and Ren Ng. NeRF: Representing scenes as neural radiance fields for view synthesis. In *ECCV*, 2020. 1
- [3] Johannes Lutz Schönberger and Jan-Michael Frahm. Structure-from-motion revisited. In *CVPR*, 2016. 1
- [4] Johannes Lutz Schönberger, Enliang Zheng, Marc Pollefeys, and Jan-Michael Frahm. Pixelwise view selection for unstructured multi-view stereo. In *ECCV*, 2016. 1
- [5] Yifan Wang, Jianjun Zhou, Haoyi Zhu, Wenzheng Chang, Yang Zhou, Zizun Li, Junyi Chen, Jiangmiao Pang, Chunhua Shen, and Tong He. π^3 : Permutation-equivariant visual geometry learning. In *ICLR*, 2026. 1